

Introduction to Neural Networks


Deep learning architectures—such as **Recurrent Neural Networks (RNNs)**, **Convolutional Neural Networks (CNNs)**, and **Deep Belief Networks (DBNs)**—have been widely applied across various domains, including *computer vision, speech recognition, natural language processing, machine translation, bioinformatics, drug discovery, and medical image analysis*.

What sets deep learning apart is the **depth of computation**, which enables these models to learn and represent **complex and hierarchical patterns**—a key requirement for tackling the most challenging real-world datasets.

Deep learning models are built upon the foundation of **Artificial Neural Networks (ANNs)**, which are inspired by the **structure and functioning of the human brain**.

In biology, the brain excels at solving complex problems, thanks to its network of specialized cells called **neurons** (or **nerve cells**). These neurons are the **fundamental building blocks** of the brain.

Some neurons are responsible for **receiving sensory input** from the external environment, while others **transmit motor commands** to control muscle movements.



Each neuron consists of four main parts: the **soma** (cell body), **dendrites**, **axon**, and **synapses**.

- **Dendrites:** These branch-like structures receive signals from multiple neighboring neurons.

- **Soma:** Also known as the cell body, it collects and integrates the incoming signals from the dendrites.
- **Axon:** Responsible for transmitting electrical signals away from the neuron's cell body.
- **Axon Terminals (Synapses):** These are the endpoints of the axon. They form connections with other neurons and are responsible for transferring information to those neurons through chemical or electrical signals.
- If the combined signal received by the soma is strong enough, it triggers an electrical impulse.
- This impulse is then transmitted through the axon to other neurons via synapses.

Although **biological brains** are capable of solving highly complex problems, each **neuron** contributes by handling only a **small part** of the overall task.


Similarly, in **artificial neural networks**, the system is composed of many units called **artificial neurons**, which work together to achieve a desired output.

Each **artificial neuron** performs only a **simple computation**, but when combined in large numbers, they can model and solve very complex problems.

Artificial Neuron (Linear Unit)

An **Artificial Neuron (Linear Unit)** takes an input, denoted by x .

- Each input is connected to the neuron with an associated **weight**, represented by w .
- When the input flows through this connection, it is **multiplied by the weight**.
- So, the value that reaches the neuron is $w \times x$.
- A **neural network learns** by adjusting these weights based on the error in its predictions.




The **bias** (denoted as **b**) is a special type of **weight** in a neural network.

- Unlike other weights, the **bias** is not associated with any input data.
- To represent it in diagrams, we typically use a constant input of **1**, so the value reaching the neuron from the bias is simply **b** (since $1 \times b = b$).
- The bias allows the neuron to **adjust the output independently** of the inputs, enhancing the flexibility of the network.

The value produced by the neuron is called the **output**, denoted as **y**.

- The neuron computes this by **summing all the signals** it receives through its connections.
- For a single input, the output is calculated as: $y = w \times x + b$.
- The **bias term (b)** allows the function to **shift left or right**, enhancing flexibility.
- Thus, an **artificial neuron** acts as a **mathematical function** that transforms input into output.



In artificial neural networks, the **connections between neurons** are represented by **weights**, simulating the behavior of biological neurons.

- **Positive weights** represent **excitatory connections**, which promote activation.

- **Negative weights** represent **inhibitory connections**, which suppress activation.
- Each input is **multiplied by its weight**, and the results are **summed together**.
- This process is referred to as a **linear combination**.


Linear Unit as a Model

The **Linear Unit** can be used to model simple relationships. For example, let's consider a case where:

- **Input:** Height (in cm)
- **Output:** Weight (in kg)

After training the model, suppose we get the following parameters:

- **Weight (w):** 0.35
- **Bias (b):** 2



$$y = wx + b = 160 * 0.35 + 2 = 58$$

To estimate the weight of a person who is **160 cm** tall:

$$\text{Output} = w \times \text{input} + b = 0.35 \times 160 + 2 = 58 \text{ kg}$$


So, the predicted weight is **58 kg**.

Multiple Inputs in a Linear Unit

Suppose the dataset includes an additional feature, such as **gender**. We can simply add another input connection to the neuron.

- Each input is multiplied by its corresponding **weight**.
- All the weighted inputs are then **added together** to produce the output.

This approach allows the model to use **multiple features** for prediction.




$$y = w_0x_0 + w_1x_1 + b$$

Suppose the dataset includes three input features: **height**, **gender**, and **country**. We can add one input connection to the neuron for each feature.

- Each feature is multiplied by its corresponding **weight**.
- All the weighted inputs are **summed** to compute the final output.

This allows the neuron to process **multiple (three) input features** in the prediction task.



$$y = w_0x_0 + w_1x_1 + w_2x_2 + b$$


Perceptron

The **Perceptron** was developed in **1958** by **Frank Rosenblatt**, based on earlier research by **Warren McCulloch** and **Walter Pitts**.

- A perceptron is an **artificial neuron** that uses the **Heaviside step function** as its activation function.
- It is used in **supervised learning** to train **binary classifiers**.
- A binary classifier is a function that determines whether a given input vector **belongs to a specific class**.

Activation Function

The **activation function** is applied to the **weighted sum of the inputs** to produce the neuron's output.




- It maps the weighted sum to an output value, usually in the range of **0 to 1**.
- It determines whether the neuron should be **activated** or not, hence the name.

How a Perceptron Works

In this example, the **perceptron** has three inputs: x_0 , x_1 , and x_2 . In general, a perceptron can accept any number of inputs.

To determine the output, **Frank Rosenblatt** introduced a simple rule:

- Each input x_i is assigned a weight w_i , representing the importance of that input.
- The perceptron computes the **weighted sum**: $\sum_i w_i \times x_i$
- This sum is compared to a **predefined threshold** (a real number).



The output is then calculated as:

$$f(x) = \begin{cases} 0 & \text{if } \sum_i w_i \cdot x_i \leq \text{threshold} \\ 1 & \text{if } \sum_i w_i \cdot x_i > \text{threshold} \end{cases}$$

In a perceptron, the **bias** (b) can be defined as the **negative of the threshold**:

$$b \equiv -\text{threshold}$$

This reformulation simplifies the perceptron's computation, allowing us to express the decision rule without explicitly using the threshold.


$$f(x) = \begin{cases} 0 & \text{if } w \cdot x + b \leq 0 \\ 1 & \text{if } w \cdot x + b > 0 \end{cases}$$

That's all there is to the basic working of a perceptron!

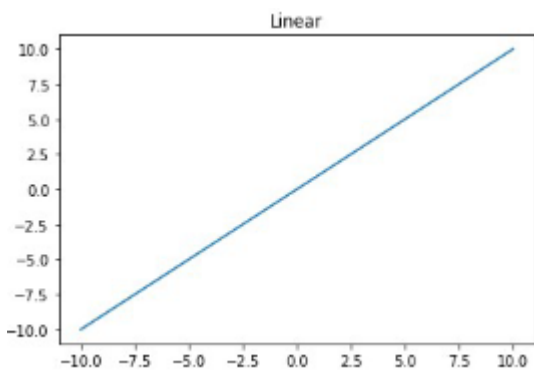
When a perceptron has **two inputs**, its input space can be visualized as a **2D plane**.

- Each input represents a point on this 2D plane.

- The neuron forms a **decision boundary**, which appears as a **straight line** (also called a *subspace* or *hyperplane* in higher dimensions) dividing the space into different classes.




$$w_0x_0 + w_1x_1 + b = 0$$




When a perceptron has **three inputs**, the input space becomes a **3D space**.

- Each input corresponds to a point in this 3D space.
- The neuron creates a **decision boundary** that can be visualized as a **2D plane**—also referred to as a *hyperplane* or *subspace*—which separates different regions in the 3D input space.




$$w_0x_0 + w_1x_1 + w_2x_2 + b = 0$$



Solving Logical OR Using a Linear Perceptron

OR Truth Table

Input x_0	Input x_1	OR Output
0	0	0
0	1	1
1	0	1
1	1	1



Perceptron Function

The perceptron computes output using:

$$y = 1 \text{ if } (w_0x_0 + w_1x_1 + b > 0), \text{ else } 0$$

Choose Weights and Bias

- $w_0 = 0.5$
- $w_1 = 0.5$
- $b = -0.25$

The function becomes: $y = 1 \text{ if } (0.5 * x_0 + 0.5 * x_1 - 0.25 > 0)$

Test with Inputs

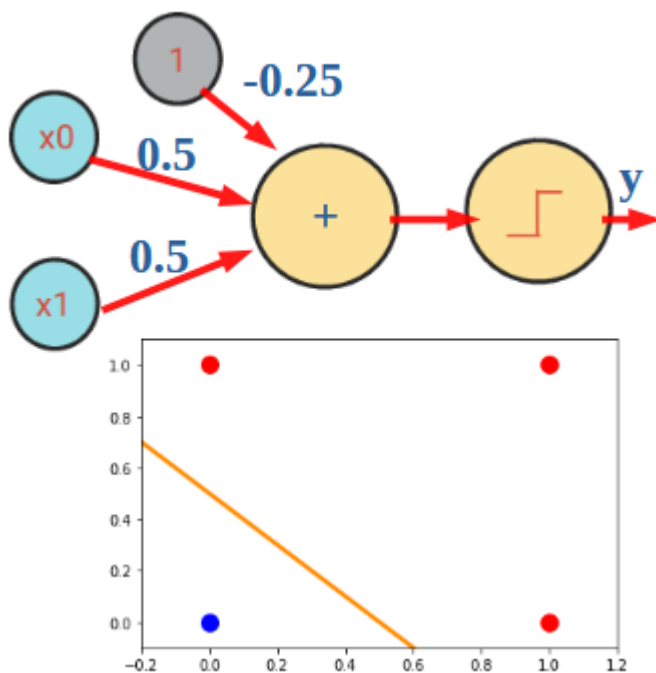
x_0	x_1	$0.5x_0 + 0.5x_1 - 0.25$	Output
0	0	-0.25	0
0	1	0.25	1
1	0	0.25	1
1	1	0.75	1

Geometric View

In the 2D input space, the perceptron forms a linear decision boundary defined by:

$$0.5x_0 + 0.5x_1 - 0.25 = 0$$

Simplified as: $x_0 + x_1 = 0.5$



This line separates the input $(0, 0)$, which produces output 0, from the others like $(0, 1)$, $(1, 0)$, and $(1, 1)$ which lie above the line and produce output 1.

Why It Works?

- Logical OR is a **linearly separable function**.
- A single-layer perceptron is sufficient to classify OR correctly.


Multi Layer Perceptron

Single-layer perceptron can solve linearly separable problems like Logical OR, but not nonlinear problems like XOR (Exclusive OR).

XOR Truth Table

Input x_0	Input x_1	XOR Output
0	0	0
0	1	1
1	0	1

1	1	0
---	---	---



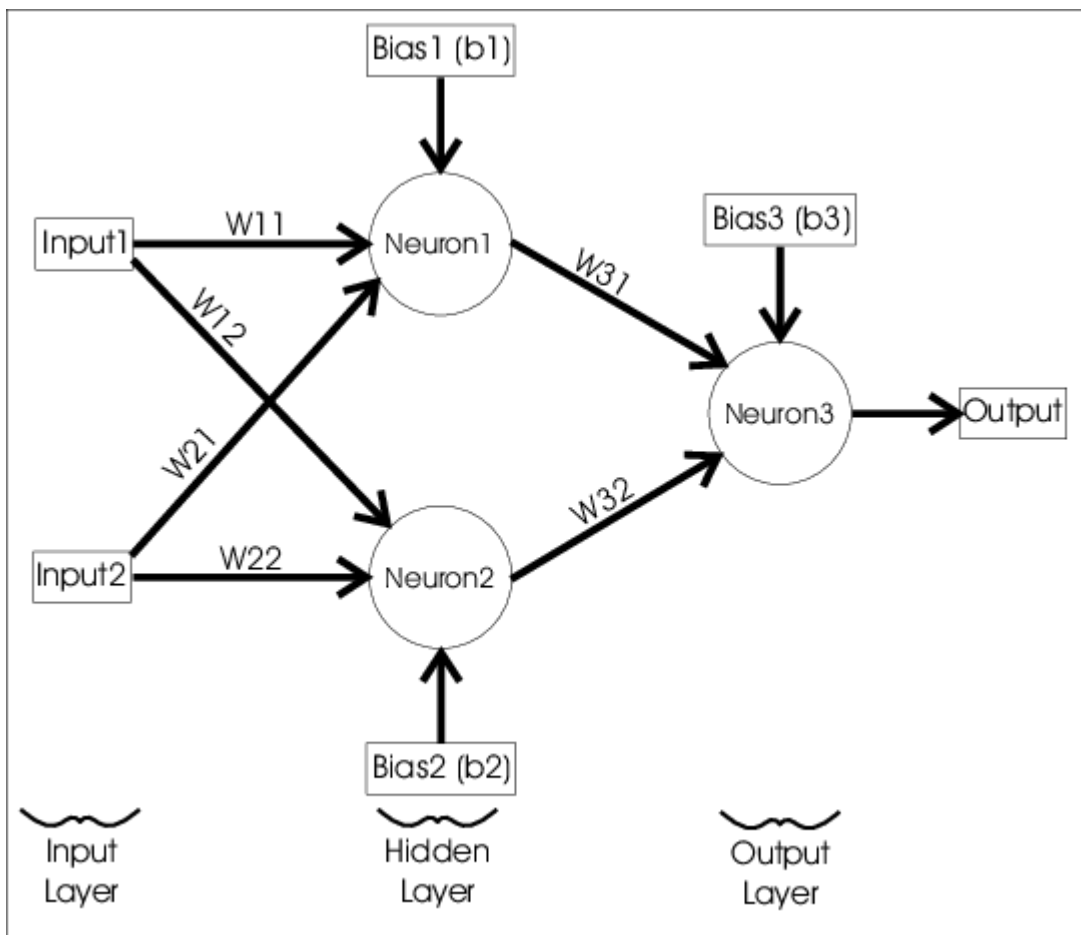
There's no straight line that can separate the 1s from the 0s — this makes XOR non-linearly separable. Therefore, a linear decision boundary created by a single-layer perceptron won't work.

We can combine multiple perceptrons in layers to learn non-linear decision boundaries. This setup is called a **Multi-Layer Perceptron (MLP)**.

Network Architecture

- **2 input neurons** (x_0, x_1)
- **1 hidden layer** with 2 neurons
- **1 output neuron**

This architecture allows the network to model non-linear functions like XOR, which are not linearly separable.



The hidden layer learns intermediate representations that transform the input space into a form where a linear classifier (like the output neuron) can separate the classes correctly.

Inputs

Let's denote the input vector as:

$$\mathbf{x} = [\mathbf{x}_0 \ \mathbf{x}_1]$$

First Layer (Hidden Layer)

- **Weights:**

$$\mathbf{W} = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$$

- **Biases:**

$$\mathbf{c} = \begin{bmatrix} 0 \\ -1 \end{bmatrix}$$

Second Layer (Output Layer)

- **Weights:**

$$w = [1 \quad -2]$$

- **Bias:**

$$b = 0$$

In hidden layers let us use step activation function

Step Activation Function

$$\text{step}(z) = \{ 1 \text{ if } z > 0, 0 \text{ otherwise } \}$$

1. Input: (0, 0)

Hidden Layer Output:

$$\begin{aligned} z &= W \cdot x + c = \\ &\begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} \cdot \begin{bmatrix} 0 \\ 0 \end{bmatrix} + \begin{bmatrix} 0 \\ -1 \end{bmatrix} \\ &= \begin{bmatrix} 0 \\ -1 \end{bmatrix} \end{aligned}$$

$$\text{step}(z) = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Output Layer:

$$y = w \cdot h + b = [1 \quad -2] \cdot \begin{bmatrix} 0 \\ 0 \end{bmatrix} = 0$$

$$\text{Output} = 0$$

2. Input: (0, 1)

Hidden Layer Output:

$$\begin{aligned} z &= W \cdot x + c = \\ &\begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} \cdot \begin{bmatrix} 0 \\ 1 \end{bmatrix} + \begin{bmatrix} 0 \\ -1 \end{bmatrix} \\ &= \begin{bmatrix} 1 \\ 0 \end{bmatrix} \end{aligned}$$

$$\text{step}(z) = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

Output Layer:

$$y = w \cdot h + b = \begin{bmatrix} 1 & -2 \end{bmatrix} \cdot \begin{bmatrix} 1 \\ 0 \end{bmatrix} = 1$$

$$\text{Output} = 1$$

3. Input: (1, 0)

Hidden Layer Output:

$$\begin{aligned} z &= W \cdot x + c = \\ &\begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} \cdot \begin{bmatrix} 1 \\ 0 \end{bmatrix} + \begin{bmatrix} 0 \\ -1 \end{bmatrix} \\ &= \begin{bmatrix} 1 \\ 0 \end{bmatrix} \end{aligned}$$

$$\text{step}(z) = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

Output Layer:

$$y = w \cdot h + b = \begin{bmatrix} 1 & -2 \end{bmatrix} \cdot \begin{bmatrix} 1 \\ 0 \end{bmatrix} = 1$$

Output = 1

4. Input: (1, 1)

Hidden Layer Output:

$$\begin{aligned} z &= W \cdot x + c = \\ &\begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} \cdot \begin{bmatrix} 1 \\ 1 \end{bmatrix} + \begin{bmatrix} 0 \\ -1 \end{bmatrix} \\ &= \begin{bmatrix} 2 \\ 1 \end{bmatrix} \end{aligned}$$

$$\text{step}(z) = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

Output Layer:

$$y = w \cdot h + b = \begin{bmatrix} 1 & -2 \end{bmatrix} \cdot \begin{bmatrix} 1 \\ 1 \end{bmatrix} = -1$$

Output = 0

Final Output (XOR Logic):

x ₀	x ₁	Output
0	0	0
0	1	1
1	0	1
1	1	0